

What is Machine Learning?

Oxford Spring School in Advanced Research Methods

Dr Thomas Robinson, LSE

Day 1 of 5

Introduction

This course

5 × 3 hour sessions, covering:

1. Introduction to ML (maximum likelihood estimation)
2. Regularised methods (bias variance trade-off)
3. Tree-based methods (hyperparameter tuning)
4. Neural networks (feature engineering)
5. Ensemble methods

The logic:

- ▶ Understand the underlying mechanics of parameter estimation
- ▶ Start from a modeling strategy we are familiar with...
- ▶ ... and move on to algorithmically more complicated cases
- ▶ Build on the same foundational concepts across each day

What won't we cover?

- ▶ ML is a broad and contested domain
- ▶ We will not cover some topics:
 - ▶ Unsupervised methods
 - ▶ Clustering algorithms
 - ▶ Most text-specific ML models in depth
 - ▶ Foundation models / large language models in depth

This course prioritises:

- ▶ Intuitive understanding of popular ML methods
- ▶ Transferability of fundamentals
- ▶ Relevance to “downstream” social science research problems

Session structure

Each day will be a mixture of:

- ▶ Lecture content (approx. 2 hours)
 - ▶ Plenty of opportunities for questions
 - ▶ Time for us to think through some applied problems
- ▶ Coding walkthroughs (approx. 1 hour)
 - ▶ Conducted in R using RStudio
 - ▶ Workflow ideas transfer directly to Python / scikit-learn style tooling
 - ▶ Hands on experience using different algorithms
- ▶ Self-guided learning
 - ▶ Applied readings using the methods we discuss in class
 - ▶ Completing and consolidating coding exercises

Slides and code are available in the shared course repository

Today's session

Goals:

1. Introduce the topic of machine learning
 - ▶ What is ML?
 - ▶ How do we distinguish ML from “traditional” statistics?
 - ▶ What sort of problems might we apply it to?
2. Introduce key conceptual distinctions we will make throughout the course
3. Introduce maximum likelihood estimation
 - ▶ A key way in which ML parameters are estimated
 - ▶ Build our own logistic regression estimator
4. Introduce a modern ML workflow for research
 - ▶ Benchmarks, held-out data, and leakage
 - ▶ Why “good prediction” is not the same as “good research design”

What is machine learning?

(Machine) learning and statistics

*“There are two cultures in the use of statistical modeling to reach **conclusions from data**. One assumes that the data are generated by a given stochastic data model. The other uses algorithmic models and treats the data mechanism as unknown” – Breiman 2001*

*“Statistical learning refers to a set of tools for modeling and understanding **complex** datasets. It is a recently developed area in statistics and blends with parallel developments in **computer science** and, in particular, machine learning” – James et al. 2013*

*“Machine learning is a subfield of **computer science** that is concerned with building algorithms which, to be useful, rely on a collection of examples of some phenomenon. . . ”
– Burkov 2019*

Machine learning

Expectation: I need a \$1m super computer

Reality: It runs in *minutes* seconds on a *personal computer* smartphone



Why machine learning?

Machine learning can be:

- ▶ Powerful
 - ▶ With respect to computational efficiency
- ▶ Flexible
 - ▶ With respect to **data generating processes**
- ▶ Reduce the burden on the researcher
 - ▶ With respect to generating data...
 - ▶ and estimating models

But ML is not a panacea!

- ▶ Cannot solve problems of poor research design
- ▶ Increased flexibility *can* lead to poor interpretability

Machine learning and social science

ML also introduces issues of its own:

Twitter apologises for 'racist' image-cropping algorithm

Users highlight examples of feature automatically focusing on white faces over black ones

Sarah Silverman sues OpenAI and Meta claiming AI training infringed copyright

A research workflow for ML in social science

Before choosing an algorithm, decide:

- ▶ What is the research task?
 - ▶ Prediction, measurement, heterogeneity, or downstream inference?
- ▶ What is the evaluation plan?
 - ▶ Train/validation/test split or resampling before model comparison
- ▶ What is the benchmark?
 - ▶ Always compare against a simple baseline model
- ▶ What are the likely failure modes?
 - ▶ Leakage, subgroup error, bad labels, and distribution shift

Running case study for the week

To keep the course coherent, we will repeatedly return to one broad applied task:

- ▶ Survey-based political prediction and measurement using CCES-style data
- ▶ Start with a simple binary classification problem
- ▶ Then ask how regularisation, trees, neural nets, and ensembles change:
 - ▶ predictive performance
 - ▶ interpretability
 - ▶ uncertainty
 - ▶ substantive usefulness

What leakage looks like in social science

Leakage means the model sees information that would not really be available at prediction time.

Common examples:

- ▶ Using post-treatment variables in a causal prediction task
- ▶ Randomly splitting panel data where the same unit appears in train and test
- ▶ Imputing or scaling using the full dataset before the split
- ▶ Predicting an outcome with text or metadata collected after the event
- ▶ Letting geography, time, or survey administration variables act as shortcuts

Leakage can make a weak research design look like a strong model.

Missing data is part of the workflow

In social science, missingness is rarely a side issue:

- ▶ Complete-case analysis can change the target population
- ▶ Missingness itself may be substantively informative
- ▶ Imputation must be done inside the resampling workflow
 - ▶ otherwise we leak information across train and test data

Later in the week we will revisit ML as a tool for imputation and measurement.

Where ML is especially useful in social science

Common high-value uses include:

- ▶ Measurement
 - ▶ Estimating ideology, topic, sentiment, trust, or state capacity
- ▶ Prediction
 - ▶ Forecasting turnout, protest, conflict, or administrative outcomes
- ▶ Heterogeneity
 - ▶ Detecting non-linearities and interactions that are hard to pre-specify
- ▶ Scaling evidence
 - ▶ Coding text, images, audio, or other high-volume observational data

Prediction and inference

The classic dichotomy when introducing ML:

$$\hat{y}_i = \hat{\beta}_0 + \hat{\beta}_1 x_{1i}$$

- ▶ **Inference:** estimating the size/direction of the (causal) relationship between variables ($\hat{\beta}$ problems)
- ▶ **Prediction:** estimating the outcome, using the relationships between variables ($\hat{y}$ problems)

These two facets are connected:

- ▶ Knowing the size/direction of (all) relationships $\rightarrow$ predict the outcome
- ▶ But we rarely know the true model
- ▶ Sometimes we can get good at $\hat{y}$ problems without knowing $\hat{\beta}$

$\hat{y}$ and $\hat{X}$ problems

We can think about where a prediction problem lies:

- ▶ $\hat{y}$ (dependent variable)
 - ▶ To predict the probability of revolution...
 - ▶ ... or the weather tomorrow
 - ▶ These are not necessarily *inferential* problems
- ▶ $\hat{X}$ (independent variables):
 - ▶ Dimensions of interest that may be important to our theory, but which are:
 - ▶ Latent (i.e. not directly observable)
 - ▶ Difficult to measure “by hand”
 - ▶ Use ML to make predictions over X
 - ▶ These estimates are used *downstream* to test a theory

Classification and prediction

Within both problems, we can distinguish two types:

- ▶ **Prediction** – estimating the value of a continuous variable (sometimes referred to as a “regression” problem) e.g.,
 - ▶ The ideology of a politician
 - ▶ The number of votes received by a candidate
- ▶ **Classification** – estimating which *class* of a category an observation belongs to, e.g.,
 - ▶ Party identity (Republican/Democrat/Libertarian/Independent)
 - ▶ The topic of a tweet (foreign/domestic/personal, positive/negative)
 - ▶ Recidivism

Supervised methods

Supervised methods (the focus of this course)

Learning the relationship:

- ▶ **Training data** are a set of labelled examples, where y is our target prediction
- ▶ We use these examples to “learn” the relationship between y and X
- ▶ Then predict y for a *new* unlabeled dataset (i.e. where the target variable is not observed)

$$\underbrace{\begin{bmatrix} 1 \\ 0 \\ 0 \\ \vdots \\ 1 \end{bmatrix}}_{\mathbf{y}^{\text{TRAIN}}} \quad \underbrace{\begin{bmatrix} 3.3 & 1.1 & 0 \\ 2.7 & 0.8 & 0 \\ 1.8 & 0.1 & 1 \\ \vdots & \vdots & \vdots \\ 5 & 1.2 & 0 \end{bmatrix}}_{\mathbf{X}^{\text{TRAIN}}}$$

Predicting on new data

$$\mathbf{X}^{\text{TEST}} = \begin{bmatrix} 3.5 & 1.9 & 1 \\ 5.4 & 0.3 & 0 \\ 1.7 & 0.5 & 1 \end{bmatrix}$$

Notation

Algebra

Throughout the course, we will follow the notation set out in Burkov (2019):

- ▶ θ and x are scalar - i.e. a single number
 - ▶ E.g., $\theta = 3.141$, $x = 1$ etc.
- ▶ $\boldsymbol{\theta}$ and $\boldsymbol{x}$ are vectors - i.e. an ordered list of scalar values
 - ▶ E.g., $\boldsymbol{\theta} = [0.5, 3, 2]$
- ▶ $\boldsymbol{\Theta}$ and $\mathbf{X}$ are matrices
 - ▶ E.g.,

$$\mathbf{X} = \begin{bmatrix} 1 & 5 \\ 24 & -3 \end{bmatrix}$$

- ▶ $\mathbf{x}^{(k)}$ is the k th column of $\mathbf{X}$
- ▶ $\mathbf{x}_i$ is the i th row in matrix $\mathbf{X}$
- ▶ $x_i^{(k)}$ is the k th element of the row vector $\mathbf{x}_i$

Probability notation

Let p denote a probability distribution *function* that returns the probability of an event or observation:

- ▶ $p(A) = 0.5$ means the probability of event A is 0.5
- ▶ All probabilities are bounded between 0 and 1
 - ▶ $0 \leq p(A) \leq 1$

Conditional probabilities means the probability of an event **given** the value of another variable

- ▶ $p(A|c) = 0.25$

Probability rules

Probabilities have some nice features:

- ▶ $p(A \text{ and } B) = p(A)p(B|A)$
- ▶ If $p(A) \perp p(B)$, then $p(B|A) = p(B)$

As a result:

- ▶ If $p(A) \perp p(B)$, then $p(A \text{ and } B) = p(A)p(B)$
- ▶ These rules explain why the probability of getting heads on both flips is 0.25:
 - ▶ $P(\text{Flip 1} = \text{Heads}) = 0.5$
 - ▶ $P(\text{Flip 2} = \text{Heads} | \text{Flip 1} = \text{Heads}) = P(\text{Flip 2} = \text{Heads})$
 - ▶ $P(\text{Flip 1} = \text{Heads} \text{ *and* Flip 2} = \text{Heads}) = 0.5 \times 0.5 = 0.25$

Notation quiz

What are the following:

1. $\mathbf{a}$
2. y_i
3. β
4. β
5. Θ

If $p(A) = 0.5$, $p(B) = 0.1$, and $p(B|A) = 0.3$:

6. Is $p(A) \perp p(B)$?
7. What is $p(A \text{ and } B)$?

Maximum Likelihood Estimation (a gentle introduction)

Searching the hypothesis space

A researcher wants to characterise an outcome as the linear combination of predictor variables:

$$\mathbf{y} = \beta_0 + \beta_1 \mathbf{x}^{(1)} + \dots + \beta_k \mathbf{x}^{(k)}$$

- ▶ We will set aside all inferential/theoretical concerns
- ▶ Focused on parsimonious description in a linear space

Since $\mathbf{X}$ is fixed (it's the data we observe), we need to find the **best** β :

- ▶ Need some way of searching across all possible values (“**hypotheses**”) and finding the one that best *fits* our data
 - ▶ Statistical learning!

Bayes Theorem (from a frequentist perspective)

$$\underbrace{P(A|B)}_{\text{Posterior}} = \frac{\overbrace{P(B|A)}^{\text{Likelihood}} \times \overbrace{P(A)}^{\text{Prior}}}{\underbrace{P(B)}_{\text{Evidence}}}$$

We can use Bayes formula to estimate the posterior probability of some parameter θ :

$$p(\theta|\mathbf{X}) \propto p(\mathbf{X}|\theta) \times p(\theta),$$

where $\mathbf{X}$ is the data.

Likelihood function

Let's suppose that we have no prior knowledge over θ , so we'll drop the prior and focus specifically on the likelihood:

$$\mathcal{L}(\theta) = p(\mathbf{X}|\theta)$$

How would we calculate this?

Likelihood function

Let's suppose that we have no prior knowledge over θ , so we'll drop the prior and focus specifically on the likelihood:

$$\mathcal{L}(\theta) = p(\mathbf{X}|\theta)$$

How would we calculate this?

$$\begin{aligned}\mathcal{L}(\theta) &= p(\mathbf{x}_1|\theta) \times p(\mathbf{x}_2|\theta) \times \dots \times p(\mathbf{x}_n|\theta) \\ &= \prod_{i=1 \dots N} p(\mathbf{x}_i|\theta)\end{aligned}$$

i.e. the product of the probability of each observations within $\mathbf{X}$, given θ .

Comparing likelihoods

Suppose we have two alternative values of θ : $\theta^{(1)}, \theta^{(2)}$. We can calculate the likelihood *ratio* (LR) of these two possible parameter values:

$$LR = \frac{\mathcal{L}(\theta^{(1)})}{\mathcal{L}(\theta^{(2)})}$$

If $LR > 1$, which set of parameters would we pick?

Comparing likelihoods

Suppose we have two alternative values of θ : $\theta^{(1)}, \theta^{(2)}$. We can calculate the likelihood *ratio* (LR) of these two possible parameter values:

$$LR = \frac{\mathcal{L}(\theta^{(1)})}{\mathcal{L}(\theta^{(2)})}$$

If $LR > 1$, which set of parameters would we pick?

► $\theta^{(1)}$

Maximum likelihood estimation

We can generalise this for all possible values of θ :

$$\arg \max_{\theta \in \Theta} \mathcal{L}(\theta) = \prod_{i=1 \dots N} p(\mathbf{x}_i | \theta)$$

i.e., from the set of all possible parameter values Θ , find the parameter value that maximises the likelihood function.

Hence, **maximum** likelihood estimation.

- ▶ How we calculate $p(\mathbf{x}_i | \theta)$ will depend on the functional form of the underlying distribution
- ▶ We'll explore this specifically with respect to logistic regression later on today

Why is this maximum likelihood a useful concept?

Numeric overflow

Multiplying many small numbers means we soon lose the power to calculate them precisely

- ▶ R double-precision numbers range from 2×10^{-308} to 2×10^{308}
- ▶ If 400 observations have $p_{\theta} = 0.01$, $\mathcal{L}(\theta)$ will be outside the computable range

What if we take the log?

- ▶ The log function is strictly increasing
- ▶ $\text{Log}(a \times b) = \text{Log}(a) + \text{Log}(b)$ so we can simply add the values

With the logged likelihood function we do not have the problem of numeric overflow!

Negative log-likelihood

We can also calculate the *negative* log-likelihood:

- ▶ I.e. put a minus sign in front!
- ▶ We then *minimise* the negative log-likelihood
- ▶ We typically want to minimise rather than maximise because many of our procedures for optimisation are based on the former
- ▶ But, broadly, this is just semantics:
 - ▶ Minimising the negative log-likelihood is the same as maximising the log-likelihood

Logistic regression

Logistic regression:

- ▶ Allows us to estimate β parameters when we have a binary outcome variable
- ▶ More broadly, it is a **binary classification** algorithm – what is the probability that $y_i = 1$ given a vector of features $\mathbf{x}_i$?

We can write the logistic regression function as,

$$f_{\theta,b}(\mathbf{X}) = \frac{1}{1 + e^{-(\theta\mathbf{X}+b)}}.$$

The goal is to find the *best* values of θ and b that “explains” the data

- ▶ For simplicity, let's include b within θ s.t. $\theta = \{b, \theta_1, \dots, \theta_k\}$

MLE of logistic regression I

For a given vector of scalar values θ , we can ask what the likelihood of the data is given those values

- ▶ With the optimal parameter choice θ^* :
 - ▶ When $y_i = 1$, $f_{\theta^*}(\mathbf{x}_i) = 1$
 - ▶ When $y_i = 0$, $f_{\theta^*}(\mathbf{x}_i) = 0$

Why can't we just use the predicted probabilities as the likelihood?

MLE of logistic regression I

For a given vector of scalar values θ , we can ask what the likelihood of the data is given those values

- ▶ With the optimal parameter choice θ^* :
 - ▶ When $y_i = 1$, $f_{\theta^*}(\mathbf{x}_i) = 1$
 - ▶ When $y_i = 0$, $f_{\theta^*}(\mathbf{x}_i) = 0$

Why can't we just use the predicted probabilities as the likelihood?

- ▶ Predicted probabilities work well when $y_i = 1$
- ▶ But we would erroneously down-weight the likelihood for all $y_i = 0$

How can we fix this?

- ▶ When $y_i = 0$, let the likelihood equal $1 - f_{\theta^*}(\mathbf{x}_i)$

MLE of logistic regression II

We construct the likelihood for *any* datapoint using a mathematical “logic gate”:

$$\mathcal{L}_{\theta} = f_{\theta}(\mathbf{X})^y \times (1 - f_{\theta}(\mathbf{X}))^{(1-y)},$$

as when $y_i = 0$, $x^{y_i} = 1$ and $x^{(1-y_i)} = x$, and *vice versa*.

Simplifying, since $f_{\theta}(\mathbf{x}_i) = \hat{y}_i$:

$$\mathcal{L}_{\theta} = \hat{\mathbf{y}}^y (1 - \hat{\mathbf{y}})^{1-y}$$

MLE optimization

We can then apply our “tricks” to make the computation easier:


$$-Log(\mathcal{L}_{\theta}) = -\sum_{i=1}^N Log(\mathcal{L}_{\theta}(\mathbf{x}_i)),$$

with the goal of minimising this quantity through choosing θ .

How exactly do we minimize this function?

- ▶ Unlike OLS it is not possible to minimize this function analytically
- ▶ We therefore have to use computation to iterate through values of θ to *approximate* the minima

Minimising the negative log-likelihood in one dimension



Gradient descent algorithm

To find the minimum of the negative log-likelihood we:

1. Choose a value for the starting parameter θ
2. Calculate the slope of the function at that point
3. Adjust our value of θ in the *opposite* direction to the slope coefficient's sign
4. Recalculate the slope, and repeat 2-4

We can generalise this to θ :

- ▶ Let $Q(\theta)$ be the negative log likelihood function
- ▶ Calculate the **gradient vector** of the function in k -dimensions
- ▶ Adjust each parameter $\theta_k \in \theta$ by the negative of the corresponding gradient element

$$\theta_k = \theta_k - \frac{\partial Q(\theta)}{\partial \theta_k}$$

Logistic regression gradient


The partial derivative for any predictor $\mathbf{x}^{(j)}$ for the *logistic* cost function is:

$$\frac{\partial Q^{\text{Logit}}}{\partial \theta_j} = (f_{\theta_j}(\mathbf{X}) - \mathbf{y}) \mathbf{x}^{(j)}$$

Hence the gradient of the function's curve for any vector of logistic parameters $\boldsymbol{\theta}$ can be described as:

$$\nabla = \begin{bmatrix} \frac{\partial Q^{\text{Logit}}(\boldsymbol{\theta})}{\partial \theta_1} \\ \frac{\partial Q^{\text{Logit}}(\boldsymbol{\theta})}{\partial \theta_2} \\ \vdots \\ \frac{\partial Q^{\text{Logit}}(\boldsymbol{\theta})}{\partial \theta_k} \end{bmatrix}$$

Progression of the descent algorithm



Learning rate

As we iteratively adjust the value of our parameter:


- ▶ It's possible we keep jumping over the minima
- ▶ Or we get stuck in a rut and the estimator fails to find an even better parameter choice

So we can scale the impact of the current gradient on the new parameter choice:

- ▶ Let's call this scalar the **learning rate** (λ)
- ▶ $\theta_{\text{New}} = \theta - \lambda \nabla$

The choice of λ is down to the researcher:

- ▶ Overly-large values will not converge
- ▶ Overly small values may take too long, or risk converging on *local* minima



Stochastic gradient descent

Gradient descent can be **expensive**:

- ▶ We have to evaluate all rows in our training data before making any updates to the parameters
- ▶ If we have lots of observations
 1. Each calculation takes a long time
 2. Takes many iterations to optimise
- ▶ Instead we can use **stochastic gradient descent** (SGD)
 - ▶ Inspect the loss of each observation (or a random subset) individually
 - ▶ Update the coefficients based on each observation

Stochastic gradient descent

Under GD, for each iteration:

$$\theta_k \leftarrow \theta_k - \lambda \sum_{i=1}^N \left(f_{\theta_k}(\mathbf{x}_i) - y_i \right) \mathbf{x}_i^{(k)}$$

Under SGD, for each iteration:

$$\theta_k \leftarrow \theta_k - \lambda \left(f_{\theta_k}(\mathbf{x}_i) - y_i \right) \mathbf{x}_i^{(k)}, \text{ for } i \in \{1, \dots, N\}$$

- ▶ SGD typically converges a lot faster than GD
 - ▶ Every iteration we make N small changes to the parameter estimate
 - ▶ Computationally more efficient (we'll cover this more later in the week)
 - ▶ At the cost of some additional noise in the optimisation process

Readings

Four suggested readings after today's class

- ▶ James, Witten, Hastie, and Tibshirani, *An Introduction to Statistical Learning*
 - ▶ Ch. 2 for the supervised learning setup
- ▶ Mullainathan and Spiess (2017), "Machine Learning: An Applied Econometric Approach"
- ▶ Grimmer, Roberts, and Stewart, *Text as Data*
 - ▶ Read as a reminder that prediction quality depends heavily on measurement quality
- ▶ Barocas, Hardt, and Narayanan, *Fairness and Machine Learning*
 - ▶ Especially useful for thinking about subgroup error and deployment risks

Coding workshop: writing our own logistic
regression classifier